

# Title: OpenAI's Transformation: From a Non-profit to a 100 Billion Valuation

## Working Paper

**Author:** Dr. Alexandra Andhov, LL.M., S.J.D.

**Short Bio:** Alexandra Andhov is an Associate Professor in Law and Technology and Corporate Law at the University of Copenhagen. Her interdisciplinary research spans corporate law, governance, capital markets, and law and technology. She actively contributes to innovative legal research and education, with her latest work exploring the intersection of computational thinking and law, highlighting the application of technology, like algorithms and blockchain, in legal professions. Founding the Copenhagen Legal/Tech Lab, she fosters a cross-disciplinary environment for innovation in legal studies. A 2019 Fulbright Scholar at Cornell Law School, Alexandra has also held visiting roles at Oxford Law School, the University of Turin, Lund University, the University of Western Australia and the Arizona State University.

**Funding Disclosure:** In acknowledgement of the financial support received for this project, this article was made possible by funding from the Independent Research Fund Denmark and the Inge Lehmann Programme under the project titled PROFIT. (Grant number: 2099-00025B). The author expresses her gratitude for the support and resources provided, which were instrumental in the completion of this work. There is no conflict of interest concerning this research.

**Teaser:** This article analyses how OpenAI used diverse corporate forms to achieve a \$100 billion valuation in less than eight years.

**Keywords:** OpenAI, corporate structure, governance, NGO

## Abstract

The transformation of OpenAI from its altruistic non-profit origins to a commercially-driven entity valued at over \$100 billion represents a significant paradigm shift in Big Tech world. Established in 2015 with the aim of developing AI for the greater good and providing open-source resources, OpenAI has since pivoted to a profit-generating model, underscored by a strategic partnership with Microsoft. This evolution has raised questions about the stewardship of data initially gathered under its non-profit model, intellectual property rights, and financial strategies to optimise tax positions. A pivotal moment in OpenAI's trajectory was the launch of ChatGPT in November 2022, which rapidly acquired 100 million users, surpassing the adoption rates of social media like TikTok and Instagram. This article delves into OpenAI's transition, examining the ethical and strategic dimensions of moving from a non-profit through a capped-profit model to a fully commercial approach, scrutinizing the role and alignment of its board members with the organisation's mission statement. Additionally, the analysis extends to the implications of OpenAI's alliance with Microsoft, probing potential conflicts with the company's original mission and values.

## 1. Introduction

OpenAI, an artificial intelligence research laboratory, was founded in December 2015 as a non-profit organisation that aimed to research artificial intelligence (AI) to discover its potential and benefits to and for the benefit of society. At this point, industry luminaries including Elon Musk, Peter Thiel, Reid Hoffman, Jessica Livingston and Infosys, among others, pledged over \$1 billion to the nonprofit to build AI that could match or improve the technology built by other tech companies, such as Microsoft or Google (Brockam, et al., 2015). The envisioned goal of OpenAI was to develop open-source software and applications that allow others to develop AI systems. The original vision was to “advance digital intelligence in the way that is most likely to benefit humanity as a whole, unconstrained by a need to generate financial return.”(Brockam, et al., 2015). A few years later, in 2019, when Microsoft agreed to invest \$1 billion, OpenAI changed the narrative and the research laboratory transitioned into a for-profit organisation ( Bass, 2019). As such, after more than three years of collecting and analysing data that others shared under the assumption that OpenAI was a non-profit organisation, this data became a vessel for developing an AI dominion. This corporate structure allowed OpenAI to bypass copyright protection laws and minimise their tax obligations.

OpenAI continues with its growth and transformation, and on November 30, 2022, it launched a prototype of ChatGPT. ChatGPT is an AI chatbot that, within two months, has reached 100 million users, breaking all the past Internet and technology records (Milmo and agency, 2023) By comparison, it took TikTok more than nine months after its global launch to reach 100 million users, whereas Instagram reached this number in more than two years. Within less than a few weeks, everyone talked about ChatGPT as the future tool to reshape all industries and working habits. OpenAI’s text-generating AI chatbot has taken the world by storm. It is able to generate various text, carry out simple analyses, and create simple code while promising hyper-charging productivity. However, as a product, it can be inadequate and make mistakes. It does not provide any references and offers a multitude of disclaimers to protect itself from any possible liability. ChatGPT was originally developed as a tool to benefit the whole of humanity, yet it ultimately benefits itself (and its founders) at the expense of humanity.

Within six months of the ChatGPT launch, a lot has happened. On February 1, 2023, OpenAI launched a subscription plan for ChatGPT Plus, starting at \$20 monthly, aiming to monetise the previously promised freely accessible AI tool. On February 7, One of OpenAI’s investors, Microsoft, integrated OpenAI’s GPT-4 model into Bing, which might have a major effect on its competition, such as Alphabet’s Google, possibly raising anti-trust concerns.<sup>1</sup> On March 14, OpenAI launched the GPT-4 model through ChatGPT Plus, allowing its paying customers to use a better model while leaving the old one for non-paying customers. On March 23, OpenAI launched plugins for ChatGPT, granting it access to third-party knowledge sources and databases, including the Internet. As OpenAI grows and adds more features to its products, the ability to steal most of Google’s 250 million unique visitors per month does not seem out of the realm of possibility. What was once a non-profit could end up becoming one of the most profitable companies of the last 20 years. Yet, how OpenAI got to such a position is indeed innovative, and it used the corporate structure and

---

<sup>1</sup> Currently, the European Commission initiated a review of agreements between tech giants and generative AI makers. See Ryan Browne, ‘Microsoft’s Multibillion-Dollar Investment in OpenAI Could Face EU Merger Probe’, *CNBC*, January 9, 2024, available online at <<https://www.cnbc.com/2024/01/09/microsoft-investment-in-openai-may-face-eu-merger-probe.html>>.

corporate law in a novel way, where while building the technological solution, it benefited the most from the not-for-profit form and then at a point of product-market-fit it transformed into a for-profit entity. This paper analyses this journey and reflects on possible legal and tax considerations of this “corporate law” innovation.

## 2. Early Years of OpenAI & Collection of First Data

Sam Altman, the CEO of OpenAI, has told the story of founding OpenAI many times (Kahn, 2023). OpenAI set out on a mission to ensure that artificial general intelligence (AGI) benefits all of humanity. The organisation's founders, including Elon Musk, Sam Altman, Greg Brockman, Ilya Sutskever, and others, envisioned a future where AGI would profoundly impact society, and they aimed to navigate its development responsibly and ethically:

*“Our goal is to advance digital intelligence in the way that is most likely to benefit humanity as a whole. We believe AI should be an extension of individual human wills, and in the spirit of liberty, as broadly and evenly distributed as possible.” (Kahn, 2023).*

This has been the continuously repeated song by the CEO and the key personnel of OpenAI.

As mentioned in the introduction, OpenAI was founded not as a traditional corporation but as a not-for-profit entity, which was an unconventional choice. To understand the reasoning behind this decision, let's take a step back and look at the events leading up to it. Prior to becoming the CEO of OpenAI, Sam Altman held the position of president at Y Combinator, a renowned Silicon Valley start-up accelerator. Altman was handpicked by Y Combinator's founder, Paul Graham, in 2014 to succeed him as the accelerator's president (Graham, 2014).

Y Combinator's primary objective has been to assist start-ups in their growth and product improvement endeavours. However, one of its most significant contributions lies in facilitating connections and fostering relationships with investors and networks in Silicon Valley.<sup>2</sup> Altman's role as Y Combinator's president granted him access to a vast pool of investors and talented individuals. By leveraging his position at Y Combinator, Altman was able to forge valuable connections that proved instrumental in the early stages of OpenAI's formation. This network of investors and talented individuals not only eased the process of securing resources but also provided expertise and support.

Based on Altman's experience from Y Combinator and the funding practices, one would assume that when founding a start-up, Altman would choose one of the corporate forms. Yet he chose a non-profit entity. To be fair, Y Combinator has also supported non-profit organisations, but only fractionally. Since 2013, it has supported 25 non-profits (out of thousands), with less than what a corporate start-up gets, specifically \$100k of funding (Walker, 2017) So, how come Altman himself chose this form of legal entity?

---

<sup>2</sup> Y Combinator has funded over 4,000 startups since 2005.

The main reason, Altman claimed, was to establish a non-profit artificial intelligence research company, unconstrained by a need to generate financial returns (Brockam, et al., 2015). Yet, when OpenAI transitioned to a for-profit entity, the reason Altman stipulated was that they were in need of further funding (Loizos, 2019). Funding is the number 1 problem of every start-up, and yet it took four years for Altman, who was the president of one of the most successful accelerators, to realise that they might encounter financial difficulties despite the \$1 billion endowment (Markoff, 2015). Forgive us if this justification seems dubious. Simultaneously, it might also be relevant to mention that in the tech start-up world, there is a recognised requirement to indicate specifically who the customer is. Start-ups spend a lot of time defining and considering an ideal customer profile (Vrionis, 2024). Therefore, claiming that a product is for everyone would be questioned and challenged by any investor.

Consequently, we should consider other plausible motives that might have prompted Altman to establish OpenAI as a non-profit legal entity rather than just developing a tool for the improvement of the entire humanity. Establishing OpenAI as a non-profit organisation provided the organisation with various advantages:

- Attracting philanthropic funding - the founders of OpenAI started the venture with a \$1 billion endowment, which they either have spent in 4 years, have not fulfilled some of the pre-agreed requirements or the promised endowment did not materialise in full;
- Tax-exempt status - OpenAI enjoyed federal and state tax-exempt status, a benefit that many NGOs in the United States qualify for (U.S. Department of State, Bureau of Democracy, Human Rights, and Labor, 2021)
- Enjoyment of other advantages, such as limited liability or separate entity status; and most importantly
- Access to billions of data under the “open” disguise.

Among the reasons mentioned above, the last one appears to be particularly significant. During its early years, OpenAI operated as a research institution with a strong focus on advancing the field of artificial intelligence. It aimed to create and disseminate OpenAI publicly committed to openness and cooperation. The organisation actively pursued collaboration with other research institutions (Lawrence, 2015). During this period, OpenAI had access to abundant data resources, although specific details regarding the origins of the data and the methods of acquisition remain limited. OpenAI disclosed that it trained the ChatGPT using Reinforcement Learning from Human Feedback (RLHF), employing similar methods to InstructGPT, but with slight differences in the data collection setup. Initially, the model was trained using supervised fine-tuning with human AI trainers engaging in conversations playing both the user and an AI assistant. The trainers used model-generated suggestions to assist in composing their responses (OpenAI, 2022). Despite this disclosure, there is no definitive answer regarding the specific sources of the data or whether any of them were subject to intellectual property considerations.

The OpenAI's GPT3 model has implemented a neural network known as a transformer (Lutkevich & Schmelzer, 2023). Transformers are deep learning algorithms now commonly used in natural language processing (NLP). The transformers are designed to process sequences of data, such as text, and learn patterns that could be used for different NLP tasks. One could also explain the architecture of ChatGPT as consisting of several layers of technology combined in one, where each of the layers performs a different task. The input layer takes the prompted text and then divides it into single tokens, such as the word ‘You’

(Stollnitz, 2023). Long and less commonly used words and expressions are divided into more tokens. The process of tokenization can be explored by entering different words on OpenAI's tokenizer page (<https://platform.openai.com/tokenizer>). These tokens are then distributed into a series of hidden layers which perform the transformation of the data into vectors, which represent the meaning of the word in a mathematical form. Afterwards, the vectors are passed through multiple layers of neural networks where their relevance and correlation to other words (tokens) used are calculated. This allows the ChatGPT to focus on specific words and phrases to better understand their relevance, which then will trigger more accurate responses to the provided prompt.

When reviewing available documentation on OpenAI websites and across various online tech websites to understand how OpenAI acquired data and which data it actually used, we arrive at only limited explanation. OpenAI discloses on its General FAQ that they were “using three primary sources of information: (1) information that is publicly available on the internet, (2) information that we license from third parties, and (3) information that our users or our human trainers provide.” (OpenAI, n.d.). Regarding the first type of information, they stipulate that they “only use available information that is freely and openly available on the Internet - for example, we do not seek information behind paywalls of from the “dark web” (OpenAI, n.d.). Publicly available information does not mean that the information is there free to grab. Even publicly available information has its author(s) and might be copyrighted - and therefore not free to use.

At the end of March 2023, the Italian Data Protection Authority, *Garante*, opened an investigation into OpenAI’s data collection practices (Vanuzzo, 2023). Even though OpenAI addressed to some extent the inquiries and now allows the users to determine whether or not it stores their data (Mukherjee & Vagnoni, 2023), neither *Garante* nor other regulators have received any detailed clarification on OpenAI’s collection of unlicensed training data. This one and elemental question continues to remain unanswered. Can OpenAI freely design a tool that takes and learns from our human content, even if available publicly, and transforms it into a product?

When prompted ChatGPT on where the data that trained it came from - it provided the following answer:

“OpenAI trained me, ChatGPT, using a large dataset collected from various sources available on the internet. This dataset includes a wide range of text from books, articles, websites, and other publicly available written material. It is important to note that while the training dataset is extensive, specific details about the sources or the inclusion/exclusion criteria for the data have not been disclosed publicly by OpenAI. Furthermore, as an AI language model, I do not have direct access to information or knowledge beyond what was included in the training data.”

*Prompted on June 13, 2023 on ChatGPT May 24 version. We know that today (January 2024), the ChatGPT is already plugged into the Internet, so it actually has access to other than just training data.*

In the realm of AI development, the prevailing paradigm suggests that more data leads to better models. Notably, OpenAI’s GPT-2 model utilised a dataset comprising 40 gigabytes of text, while GPT-3, the

foundation for ChatGPT, was trained on a massive 570 gigabytes of data (Heikkilä, 2023). The GPT-4 has been trained on more than 1 trillion parameters. OpenAI's training data, which underpins the development of various GPT models, raises important questions regarding its provenance. While the specifics regarding the origin of the data utilised in GPT models are not clearly disclosed in OpenAI's terms and policies, it leaves room for uncertainty about its sources. The lack of transparency regarding the data's origins raises concerns about potential biases, ethical considerations, data quality assurance, and also copyright.

The inability to provide specific details about the data source raises concerns about the "open" ethos that OpenAI claims to uphold, a contradiction that is particularly striking given its origins. While OpenAI, as a non-profit research institution, may not have faced intense regulatory scrutiny, it is undeniable that a company generating billions in revenue should be transparent about the origin of the data - that forms a part of today's USD 100 billion evaluation. It becomes increasingly apparent that this source of wealth should have been disclosed from the outset. Considering the accessibility of data during its early years and limited scrutiny from regulators suggest that the choice of a non-profit legal entity may have been a deliberate decision from the beginning.

In light of these observations, it becomes imperative to question the motivations and intentions behind OpenAI's non-profit formation.

### **3. Corporate Transformation in 2019, Microsoft Partnership and Capped Profit**

The year 2019 was a big year for OpenAI as it underwent a substantial corporate and financial transformation. First, in March 2019, it transitioned from a non-profit to a for-profit corporation. In its announcement, OpenAI stipulated that the transition has to take place for OpenAI to be able to attract investments (Brockam & Et al., 2015). In its public statement, OpenAI claimed that in order to maintain its "openness" and continue developing its product, it created a "hybrid" of for-profit and non-profit, which they called a "capped profit" company. Before we discuss the specifics of this "hybrid" type of company and what it actually means, it is important to explain why this transition happened in the first place.

Reflecting on the reasons stipulated by OpenAI itself, the funding was the main reason. Despite the original endowment of \$1 billion, OpenAI claimed to need more financing given that the AI sector, with Google, Meta and other companies, started to become more and more competitive. This in itself raises a number of questions - were they creating AI for the greater good or creating AI to dominate the market? These are two distinct objectives, each proposing a unique path. One could argue that they represent two entirely divergent journeys. Yet, not only OpenAI performed a corporate transformation, but within a few months, in July 2019, it announced that it had concluded an exclusive partnership with Microsoft (Brockman, 2019). It is without any doubt that these two decisions have been taken in conjunction. When exactly these decisions were taken is a question for investigative journalists or competition authorities to review the communication between OpenAI and Microsoft C-level personnel to see when the discussions about the partnership truly started.

The partnership between the two entities is substantial, given that OpenAI uses Microsoft's Azure cloud technologies to run its large-scale AI models and simulations. This partnership solidifies Microsoft's position at the forefront of AI development and ultimately means that Microsoft owns the entire infrastructure on which OpenAI is built, which makes these two entities almost inseparable. Microsoft's engagement with OpenAI delves deeper than a mere financial investment - it extends into the realm of the commercialisation of OpenAI's technologies. Under the agreement, Microsoft can integrate OpenAI's innovations to enhance its products, encompassing everything from augmenting Microsoft's chatbots to infusing sophisticated language comprehension abilities into Office applications. As such, Microsoft is launching Copilot, a generative AI assistant integrated within a host of the company's workplace apps (Spataro, 2023).

Thus, we can safely conclude that the collaboration between OpenAI and Microsoft is not just an investment but a full-fledged partnership. Which currently is also being investigated by the European Commission for the potential breach of the EU Merger Regulation (Tar, 2024). The potential avenues through which Microsoft can leverage the "open" AI technologies are limitless and virtually uncapped. From Microsoft's perspective, they can seamlessly infuse OpenAI's technologies into their products with no constraints and without limitations, and participate in the profit of OpenAI when it licenses its solutions to third parties. This ultimately means a full-on market capture. Yet not only that, but this partnership fully disregards the claim of the "capped profit" entity that OpenAI established. While OpenAI could in theory, cap their profit, they will not cap the profit of Microsoft. The market participants have already figured this out, as Microsoft recently challenged Google as Microsoft took the number one place of the world's most valuable company (Randewich, 2024).

Furthermore, consider the potential head start that Microsoft might secure over its competitors. Before OpenAI releases any new version or tool to the public, Microsoft might gain a few weeks or months of an advantage. Drawing a parallel to the automotive industry, Microsoft's potential two-month advantage equates to Mercedes boasting a lead of three to five years, considering the velocity of advancements in these sectors.

Less than a year into its partnership with Microsoft, OpenAI began to roll out commercial products leveraging the power of natural language modelling. In 2,5 years after establishing the partnership with Microsoft, OpenAI has not only grown substantially, but it also launched numerous products that became widely used. The biggest success so far has been the ChatGPT, a conversational AI tool, which achieved a million users within five days after its release in November 2022 (Buchholz, 2023). For comparison, it took 2.5 months for Instagram to achieve one million users.

Today, Open AI offers an extensive array of services, such as GPT-4, an advanced language processing AI model, Codex, Dall-E and others, available via subscription. Furthermore, OpenAI provides application programming interfaces (APIs) for developers to directly access, integrate and use OpenAI's models in their own applications (Brockam & Et al., 2015). While in 2022, according to Sam Altman, OpenAI reported a loss of \$540 million (Woo & Afrati, 2023) within 12 months, it achieved annualised revenue of \$1.6 billion in 2023 (Reuters, 2023). This ultimately gave Sam Altman the title of the CEO of the Year by Time magazine (Bajekal & Perrifo, 2023), despite the internal fights for control in OpenaAI and him being fired and reinstated within five days in November 2023.

What is even more interesting than the internal struggles, at the beginning of 2023, Sam Altman predicted that the revenue of OpenAI would reach \$1 billion in 2023 (Dastin et al., 2022). And yet, with this prediction, he signed a \$10 billion deal allowing Microsoft to own 49% of the company's for-profit entity (OpenAI, 2023). One could argue that this was done due to the unforeseeable future and the need to secure the #1 place on the market, yet selling 49% of the for-profit company might be considered a questionable business decision. As a consequence, Microsoft is the leading limited partner while having a commercial agreement as part of the one billion invested in 2019. After this deal, OpenAI secured another venture round of \$300 million with 28,7 billion valuation by the likes of Sequoia Capital, Tiger Global Management, Andreessen Horowitz, and Bedrock Capital (Singh & Lunden, 2023).

## 4. Corporate Law Perspectives on OpenAI

The majority of start-ups get incorporated as limited liability companies, which gives them a lot of possibilities to pursue their business venture while limiting corporate and personal liabilities and attracting third-party funding. The corporate design for start-ups is straightforward. Yet, the corporate structure and journey were a bit different for OpenAI. We divide the journey into three parts following the above-discussed timeline.

### 4.1. 2016 – OpenAI Inc.: Non-Profit Company

When OpenAI was established - back in 2015 - it publicly stipulated that they have been established as a non-profit entity. In their post ‘Introducing OpenAI’, the officers of OpenAI stipulated that the entity is non-profit (Brockam & et al., 2015). When reviewing publicly available data from state corporate registries in Delaware and California, we found that the first established OpenAI entity was OpenAI Inc. in Delaware on December 8th, 2015. This entity has been established as a corporation. According to the publicly available data on the website of the Department of State of Delaware, there is no stipulation of this entity as being non-profit. One month after the Delaware corporation OpenAI Inc. was established, on January 7, 2016, it established a branch in California in the form of a non-profit corporation (Cal. Corp. Code §5133). Please be aware of the nuances. Not a non-profit organisation, but a **non-profit corporation**. Even the non-profit corporations in California are governed by the Californian Corporate Code. Under the Californian Corporate Code, non-profit organisations may be formed as corporations, unincorporated associations or trusts (Cal. Corp. Code §§5000 to 5080). The chosen form was a non-profit corporation, which, however, continued to be a branch of a Delaware corporation.

Given that Californian OpenAI Inc. is a non-profit corporation, it enjoys certain aspects of this corporate form, including limited liability, that it offers its shareholders. The non-profit corporation is legally created with the filing of Articles of Incorporation duly certified by the Secretary of State (Cal. Corp. Code §5053). It is governed by its board of directors. Usually, non-profit corporations have between three to 25 directors. Directors might adopt Bylaws set forth a number of directors of the corporation, qualifications, duties, and compensation of directors and various governance powers and limitations of the board of directors. Additionally, some non-profit organisations may be exempt from federal income and excise taxes and California income, franchise, property, or sales taxes. According to OpenAI’s structure, OpenAI, Inc. in California is registered as 501(c)(3) (OpenAI, 2023), which means that the non-profit holds a tax-exempt

status. Furthermore, a nonprofit corporation may also be recognised as a qualified recipient for charitable contributions deductible by the donor or the donor's estate for federal income, gift, or estate tax purposes and California income tax purposes (CEB 11th ed.). This naturally incentivises the contributors (investors) to donate to the purpose. The question is whether there is any form of consideration that they get for their contribution. This, however, is left for another investigation. Ultimately, OpenAI's choice to establish itself as a non-profit corporation not only provided a flexible operations framework and limited reporting obligations but also availed various tax benefits. The review of OpenAI's tax submissions is beyond the scope of this piece but, nonetheless, should be reviewed.

Upon examining publicly accessible data, we discovered another intriguing detail. Contrary to expectations, Samuel Altman does not serve as an active director of the Californian non-profit corporation OpenAI Inc., which is touted as the nucleus of the organisation. Instead, a corporate service company fills the roles of directors and officers. This arrangement is typically chosen for entities within corporate structures that exist in name but do not necessarily actively engage in significant tasks or operations.

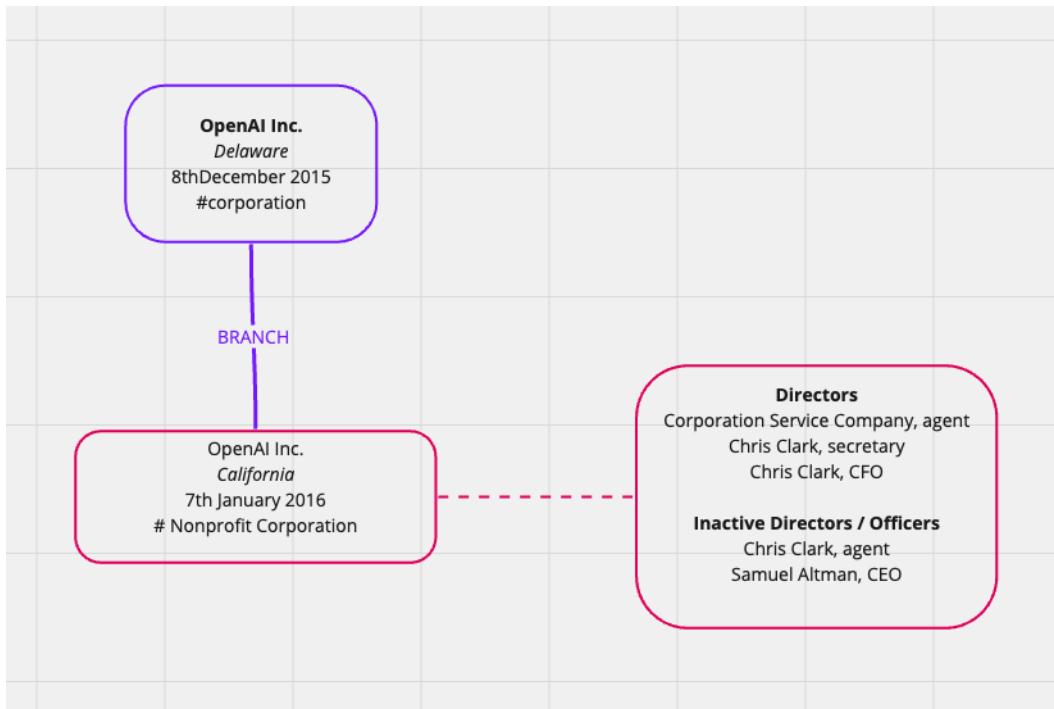


Figure 1 OpenAI Inc Structure showcasing the ownership and management of the OpenAI Inc (Data accessed through Opencorporates.com; lastly reviewed on January 19, 2024)

#### 4.2. 2018 - OpenAI, L.P.: Capped-Profit Company

In 2018, a few months before Microsoft invested \$1 billion, Open AI went through a corporate restructuring, where a new “capped-profit” company was introduced - **OpenAI, L.P** (Brockman & Sutskever, 2019). According to the announcement, the idea behind OpenAI, L.P. (limited partnership, not limited profit) is that investors and employees can get a capped return if “we” succeed at our mission (Brockman & Sutskever, 2019). The announcement does not further elaborate on who is “we”, and neither does it detail

the success factors of “our mission”. As to the corporate structure, the disclosed information in the announcement is generic, most presumably drafted with the help of one of the leading corporate firms that specialize in writing disclosures and thus leave very small space for error and liability. Following the announcement, OpenAI, L.P. - will be a for-profit Delaware Limited Partnership managed by its General Partner, a single-member Delaware LLC controlled by OpenAI, Inc. (the Nonprofit)’s Board of Directors.

According to the statement published by Greg Brockman, Ilya Sutskever and OpenAI, the structure of OpenAI, L.P. should look in the following way:

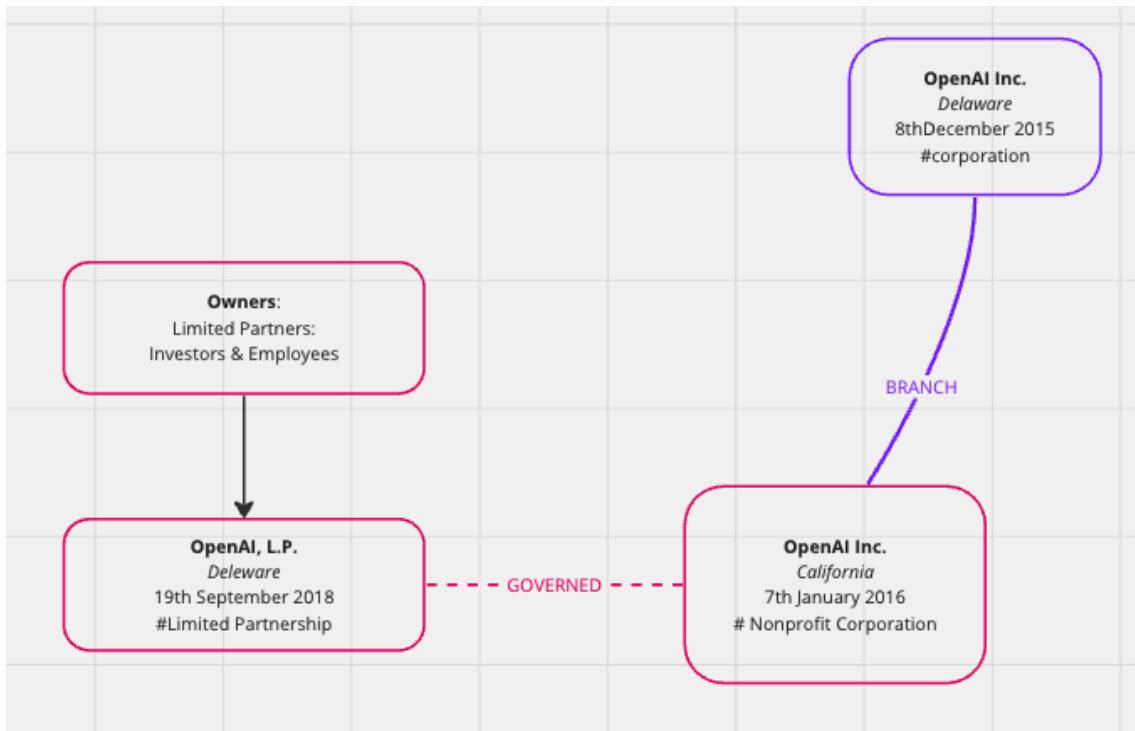


Figure 2 OpenAI, L.P. Structure According to OpenAI L.P. (Data accessed through Opencorporates.com and Delaware Department of State Registry; lastly reviewed on January 19, 2024)

When reviewing the data publicly available on OpenCorporates and at the Delaware Department of State, the OpenAI, L.P., established on 19th September 2018, with file number 7063675, is currently in the Delaware Department of State: Division of Corporations registered as OpenAI OPCO, LLC, a limited liability company. OpenAI OPCO, LLC is the mother company of OpenAI, L.P., established in California, which means that the structure looks more like:

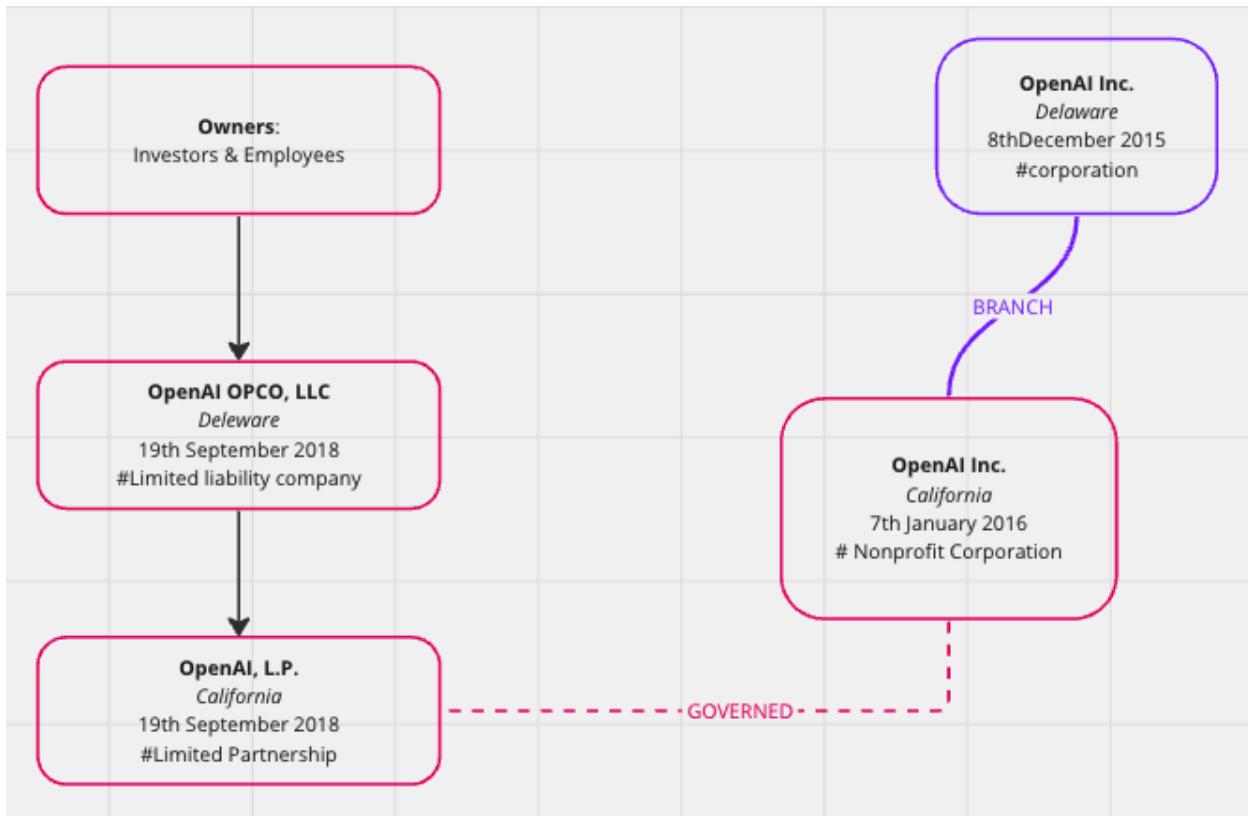


Figure 3 OpenAI, L.P. Structure (Data accessed through Opencorporates.com and <https://icis.corp.delaware.gov/>; lastly reviewed on January 30, 2024)

Drawing upon the data collected on the Opencorporate and the Delaware Department of State, it is evident that the L.P. in question was constituted under the legal framework of California rather than Delaware while having a Delaware parent company, contrary to the assertion made by OpenAI (Brockman & Sutskever, 2019). This distinction might be substantial from a number of perspectives, including fiduciary duties, tax obligations or others. California law steadfastly maintains that the fiduciary responsibilities of a general partner are immutable and cannot be diminished through any contractual arrangements (Stull v. Fox, 2010; Stull v. Fox, 2012) Conversely, Delaware law offers a degree of flexibility, permitting partners to tailor their fiduciary duties through mutual agreement (In re SunEdison, Inc., 2022). However, it is crucial to note the disparity in shareholder rights between the two jurisdictions. In California, the scope for shareholders, which in the context of OpenAI, L.P., encapsulates investors and employees to inspect internal corporate documents, is significantly narrower than in Delaware. This stark difference could provide the rationale behind the choice of law. Nonetheless, it is imperative to approach further assertions with caution, as a comprehensive analysis is impossible without complete access to the corporate and taxation documents. Any extrapolation beyond the available data would border with speculation.

More pressing, however, is the need to confront the *narrative* and the legal inconsistencies that OpenAI purports. Central to this discourse is the establishment of OpenAI, L.P. as a "capped-profit" entity, a concept that notably, lacks legal standing and recognition.

Specifically, California's corporate law, the governing statute for Californian OpenAI, L.P., does not recognize the notion of a "capped-profit" entity. According to the 2008 California Uniform Limited

Partnership Act (CA Corp Code),<sup>3</sup> any limited partnership must adhere to established legal frameworks. These partnerships offer significant advantages, including limited liability for investors like Microsoft and the benefit of pass-through taxation, allowing profits or losses to directly affect individual partners' tax liabilities. Therefore, in case the OpenAI, L.P. had any losses, the partners can sign off on these losses. The question, therefore, is, how could the partners incorporate the capped-profit nature? According to the CA Corp Code, the partners of an L.P. can agree on a "partnership agreement", a document delineating the rights, responsibilities, and power dynamics amongst various partner classes.<sup>4</sup> Thus, it is within this partnership agreement that the partners can stipulate the "capped nature" of the L.P., which currently is agreed on 100x (Brockman & Sutskever, 2019). However, this framework inherently possesses a malleable nature; partners retain the prerogative to either escalate the cap or dissolve it altogether.

If OpenAI genuinely aspired to create an entity dedicated to leveraging its profits for continuous societal betterment and ethical AI utilisation, alternative corporate structures were available, notably the California Benefit Corporation or California Social Purpose Corporation (Clark JR & Babson, 2012; Kelley, 2014). These options intrinsically align with philanthropic and societal objectives. In stark contrast, OpenAI's choice to adopt a limited partnership model—a structure predominantly favoured for corporate investments and extensively utilised by private equity and venture capital firms—casts doubt on the sincerity of their publicly declared commitment to prioritise "humanity" in their fiduciary responsibilities (OpenAI, 2018). OpenAI's proclaimed dedication to placing "humanity" at the forefront of their fiduciary duties sharply contrasts with the selected corporate framework. This disparity between stated intentions and actual corporate organisational choices highlights a significant misalignment that should be further discussed, even more so by the US and EU regulators.

### 4.3. Fictitious Governance

In November 2023, OpenAI's hybrid corporate structure became the focal point of extensive discussions when a series of high-profile events unfolded. These events centred around Samuel Altman, the co-founder and CEO of OpenAI, who experienced a dramatic ousting and subsequent reinstatement. The narrative reached its peak in the week leading up to the US Thanksgiving, marked by a decision from the Board of OpenAI Inc. (a non-profit entity detailed in section 4.1) to dismiss Altman, as reported by Metz (2023). Amidst the tumultuous sequence of events—including Altman's interim shift to Microsoft, escalating anxieties among venture capital investors, and threats of resignation from key employees—the most consequential and definitive outcome emerged: the Board itself was ultimately dismissed.

The Board, including Ilya Sutskever, also a co-founder of Open AI and its chief scientist, was concerned about the pace of OpenAI's commercialization vis-a-vis the necessary safeguards to protect against new advanced models (Varanasi, 2023). Unfortunately, no specifics were publicly announced, but on a number of occasions, the members of the Board stipulated that Altman did not fully disclose all the necessary information (Goswami & Capoot, 2023).

When OpenAI Inc was established back in 2016, the promise of the organisation and the people behind it, including its co-founders, was that this "independent" non-profit corporation's board would oversee the activities of the whole OpenAI corporate group and make sure that the vision to "advance digital

---

<sup>3</sup> This forms part of the 2022 California Corporation Code, Title 2 (15900 - 15912.07).

<sup>4</sup> CA Corp Code §15903.07. (2022).

intelligence in the way that is most likely to benefit humanity as a whole, unconstrained by a need to generate a financial return.” would be fulfilled (Brockam, et al., 2015). Yet - this was not the case. Within a week after Altman was removed, it was the majority of the Board that was removed. According to the bylaws, the Board had the competence to remove Altman, not the other way around.

In accordance with the [bylaws](#) set forth during the establishment of OpenAI Inc., the established governance model wields considerable power to the Board, which is entrusted with oversight of the OpenAI’s direction. It stipulates in Article VI, section 3 that “Subject to the rights, if any, of an officer under any contract of employment, may be removed, with or without cause, by the Board of Directors or by an officer on whom such power of removal may be conferred by the Board of Directors.” There is no such power given to the CEO. Thus, ultimately, the adopted bylaws and governance procedures were blatantly disregarded.

An additional component of this analysis involves contrasting the Board's composition pre- and post-reinstatement of Samuel Altman. This comparison underscores the shift in the company's direction and intention, reinforcing the argument that OpenAI has strayed significantly from its foundational mission and objectives. It highlights a concerning departure from the fiduciary responsibilities that were initially established, calling into question the current OpenAI's and its Board's adherence to the company's core principles and duties.

The Board that dismissed Altman was relatively diverse in expertise and perspectives. It consisted of Ilya Sutskever, OpenAI's co-founder and chief scientist, Helen Toner, a researcher focusing on the safe development of AI who had previously expressed criticism of OpenAI's launch of ChatGPT (Imbrie et al, 2023); Tasha McCauley, a senior management scientist in Rand Corporation; and Adam D'Angelo, the CEO of Quora and former CTO of Meta. Notably, Adam D'Angelo retained his seat as he reportedly negotiated Altman's return as a CEO (Goswami & Capoot, 2023).

Currently, the new Board comprises only three members. Alongside Adam D'Angelo, the Board now includes Bret Taylor, assuming the role of chair of the Board and bringing his experience as the former co-CEO of Salesforce, and Larry Summers, an economist and a former Treasury secretary during the Clinton Administration. Upon a closer examination of the members' backgrounds and specialisations, we observe a clear shift in focus. The original Board was characterised by a blend of scientific acumen and a critical perspective on technological progression. Whereas the current composition underscores a clear pivot towards growth and commercialisation. This shift implies that the Board initially envisioned as a guardian of AI's safe and beneficial development for humanity, appears now to be primarily driven by financial considerations. Moreover, it is worth noting the lack of diversity in the current Board, which is exclusively composed of white men, marking a departure from the broader representation seen previously.

Thus, the key question is whether the opaque corporate structure combined with a clear shift in governance could constitute a blatant misrepresentation.

## 5. Conclusion: New Start-up Playbook?

OpenAI's transformation from a not-for-profit organisation to one of the Big Tech companies with a valuation of \$100 billion may be charting new playbook for companies. The ethical implications of this transition cast a huge shadow over the ethos that the founders and key employees have been trying to create. The organization's pivot leveraged its non-profit status to accrue data and enjoy tax benefits, which should open hundreds of questions over the integrity of its foundational pledge to prioritise humanity's interests.

The once presumed vigilant guardians of this pledge - the Board - has transformed as well and any protection of ethical standards seem to be greatly forgotten, with only one task - further commercialization.

OpenAI's narrative could indeed inspire startups, yet it also casts a shadow on the societal trust placed in non-profit organizations. It raises a crucial question: if the public cannot place their trust in non-profit entities, who then is the custodian of societal good in the landscape of data protection? This reflection is not only about the integrity of OpenAI but also about the responsibility of the corporate world at large to uphold the values of transparency and trust that society rightfully expects.

As highlighted by Professor Carissa Véliz, the extraction of data by corporations is a silent process, often unbeknownst to individuals, causing an invisible yet profound impact (Véliz, 2020). OpenAI's shift to a profit-making entity, after harnessing vast amounts of data under a non-profit facade, signifies more than a change in business strategy; it represents a fundamental alteration in the trust dynamic between the organization and its community of users and data providers.

It is highly plausible that the sophistication with which OpenAI navigated corporate structures and governance to amass data was not accidental but rather calculated. This orchestration could be perceived as a breach of ethical standards (and possibly legal), given the potential misrepresentation to initial contributors and the multitude of users who have provided their data for not-profit-driven objectives, they were led to believe. The scarcity of publicly accessible documentation, due to OpenAI's status as a private company, and the absence of internal whistleblowers, leaves a gap in transparency that only time might fill. Though the emerging investigations by various data protection authorities, including the European Commission and the UK's Competition and Markets Authority, suggest an impending era of accountability and possibly disclosure of more information to the public that would challenge the protective veils of private corporate structures.

The interplay between the corporate structure and governance and OpenAI's data collection practices is not accidental. The decisions on structuring and re-structuring the company were carried out meticulously. Even if the corporate law does not prohibit changes in corporate structures, there is an argument to be made regarding potential misrepresentation. Misrepresentation to the organisations that in the early years of OpenAI's existence, voluntarily shared their information. Misrepresentations to the hundreds of millions of users who pay and use ChatGPT and who, on daily basis share their data.

## References

- Bajekal, N., & Perrifo, B. (2023, December 6). *Sam Altman is TIME's 2023 CEO of the year*. TIME. <https://time.com/6342827/ceo-of-the-year-2023-sam-altman/>.
- Bass, D. (2019). *Microsoft to invest \$1 billion in OpenAI*. Los Angeles Times. <https://www.latimes.com/business/story/2019-07-22/microsoft-openai>.
- Brockam, G., & et al. (2015). *Introducing OpenAI*. OpenAI. <https://openai.com/blog/introducing-openai>.
- Brockman, G. (2019). *Microsoft invests in and partners with OpenAI to support us building beneficial AGI*. OpenAI. <https://openai.com/blog/microsoft-invests-in-and-partners-with-openai>.
- Brockman, G., & Sutskever, I. (2019). *OpenAI LP*. OpenAI. <https://openai.com/blog/openai-lp>.
- Browne, R. (2024). *Microsoft's multibillion-dollar investment in OpenAI could face EU merger probe*. CNBC. <https://www.cnbc.com/2024/01/09/microsoft-investment-in-openai-may-face-eu-merger-probe.html>.
- Buchholz, K. (2023, January 24). *Infographic: Threads shoots past one million user Mark at lightning speed*. Statista. <https://www.statista.com/chart/29174/time-to-one-million-users/>.
- California Corporation Code. (2022). §§ 15900–15912.07.
- C.E.B., 1 Advising California Nonprofit Corporations 3d, § 3.1 et seq.; Cal. Transactions Forms, 5 Business Entities, Chap. 24; 9 Summary (11th), Taxation, § 192 et seq.
- Clark JR, W. J., & Babson, E. K. (2012). How Benefit Corporations are Redefining the Purpose of Business Corporations. *William Mitchell Law Review*, 38(2).
- Dastin, Krystal Hu, J., Hu, K., & Dave, P. (2022, December 15). *Exclusive: ChatGPT owner OpenAI projects \$1 billion in revenue by 2024*. Reuters. <https://www.reuters.com/business/chatgpt-owner-openai-projects-1-billion-revenue-by-2024-sources-2022-12-15/>.
- Goswami R & Cappot A, *Sam Altman's back. Here's who's on the new OpenAI board and who's out* (CNBC, November 22, 2023).
- Graham, P. (2014). *Sam Altman for president*. Y Combinator. <https://www.ycombinator.com/blog/sam-altman-for-president>.
- Heikkilä, M. (2023, April 19). *OpenAI's hunger for data is coming back to bite it*. MIT Technology Review. <https://www.technologyreview.com/2023/04/19/1071789/openais-hunger-for-data-is-coming-back-to-bite-it>
- Imbrie A, Daniels O, & Toner H, Decoding Intentions: Artificial Intelligence and Costly Signals, *Center for Security and Emerging Technology* (October 2023).
- In re SunEdison, Inc., No. 16-10992 (DSJ), 2022 WL 1223281 (Bankr. S.D.N.Y. Apr. 26, 2022).

Kahn, J. (2023). *The inside story of ChatGPT: How OpenAI founder Sam Altman built the world's hottest technology with billions from Microsoft*. Fortune. <https://fortune.com/longform/chatgpt-openai-sam-altman-microsoft/>.

Kelley, R. (2014). The Emerging Need for Hybrid Entities: Why California Should Become the Delaware of "Social Enterprise Law". *Loyola of Los Angeles Law Review*, 72, 619.

Lawrence, N. (2015). *OpenAI won't benefit humanity without data-sharing*. The Guardian. <https://www.theguardian.com/media-network/2015/dec/14/openai-benefit-humanity-data-sharing-elon-musk-peter-thiel>.

Loizos, C. (2019). *Sam Altman's LEAP of faith*. TechCrunch. <https://techcrunch.com/2019/05/18/sam-altmans-leap-of-faith/>.

Lutkevich, B., & Schmelzer, R. (2023). *What is GPT-3? Everything you need to know - TechTarget*. Enterprise AI. <https://www.techtarget.com/searchenterpriseai/definition/GPT-3>

Markoff, J. (2015). *Nytimes.com*. The New York Times - Breaking News, US News, World News and Videos. <https://www.nytimes.com/2015/12/12/science/artificial-intelligence-research-center-is-founded-by-silicon-valley-investors.html>.

Metz, C, *OpenAI's Board Pushes Out Sam Altman, Its High-Profile C.E.O.*, The New York Times (November 17, 2023).

Milmo, D. (2023). *ChatGPT reaches 100 million users two months after launch*. The Guardian. <https://www.theguardian.com/technology/2023/feb/02/chatgpt-100-million-users-open-ai-fastest-growing-app>.

Mukherjee, S., & Vagnoni, G. (2023). *Italy restores ChatGPT after OpenAI responds to regulator* Reuters, <https://www.reuters.com/technology/chatgpt-is-available-again-users-italy-spokesperson-says-2023-04-28/>.

OpenAI. (2018). *OpenAI Charter*. <https://openai.com/charter>.

OpenAI. (2022). *Introducing ChatGPT*. <https://openai.com/blog/chatgpt>.

OpenAI. (2023). *OpenAI and Microsoft extend partnership*. <https://openai.com/blog/openai-and-microsoft-extend-partnership>.

OpenAI. (2023). *Our structure*. [https://openai.com/our-structure?utm\\_source=newsletter&utm\\_medium=email&utm\\_campaign=sendto\\_newslettertest\\_business\\_&stream=top#\\_ga=2.20795512.1198283772.1705610122-1057350421.1705610122>](https://openai.com/our-structure?utm_source=newsletter&utm_medium=email&utm_campaign=sendto_newslettertest_business_&stream=top#_ga=2.20795512.1198283772.1705610122-1057350421.1705610122).

OpenAI. (n.d.). *How ChatGPT and Our Language Models Are Developed*. [help.openai.com](https://help.openai.com/en/articles/7842364-how-chatgpt-and-our-language-models-are-developed). <https://help.openai.com/en/articles/7842364-how-chatgpt-and-our-language-models-are-developed>.

Randewich, N. (2024). *Microsoft challenges Apple as the world's most valuable company*. Reuters <https://www.reuters.com/technology/microsoft-challenges-apple-worlds-most-valuable-company-2024-01-10/>.

Reuters. (2023). *OpenAI annualized revenue tops \$1.6 billion - The Information.* <https://www.reuters.com/technology/openai-annualized-revenue-tops-16-billion-information-2023-12-30/>.

Singh, J., & Lunden, I. (2023, May 4). *OpenAI closes \$300M share sale at \$27B-29B valuation.* TechCrunch. <https://techcrunch.com/2023/04/28/openai-funding-valuation-chatgpt/>.

Spataro, J. (2023, May 16). *Introducing Microsoft 365 copilot – your copilot for work.* The Official Microsoft Blog. <https://blogs.microsoft.com/blog/2023/03/16/introducing-microsoft-365-copilot-your-copilot-for-work/>.

Stollnitz, B. (2023). *How GPT models work: Accessible to everyone.* Bea Stollnitz. <https://bea.stollnitz.com/blog/how-gpt-works/>.

Stull v. Fox, 487 F. App'x 336 (9th Cir. 2012).

Stull v. Fox, No. 09-cv-6081, 2010 WL 3895538 (C.D. Cal. Oct. 4, 2010).

Tar, J. (2024, January 9). *EU commission to examine Microsoft-openai partnership.* www.euractiv.com. <https://www.euractiv.com/section/competition/news/eu-commission-to-examine-microsoft-openai-partnership/>.

U.S. Department of State, Bureau of Democracy, Human Rights, and Labor. (2021, January 19). *Non-governmental organizations (NGOs) in the United States.* <https://www.state.gov/non-governmental-organizations-ngos-in-the-united-states/>.

Vanuzzo, A. (2023). *Bloomberg.* Bloomberg - Are you a robot?. <https://www.bloomberg.com/news/articles/2023-03-31/italy-privacy-regulator-opens-probe-on-openai-over-chatgpt>.

Varanasi, L (2023) We now have more info on what Sam Altman did that was so bad he got fired from OpenAI, *Business Insider* (December 10, 2023).

Velíz, C. (2020) Privacy is Power.

Vrionis, J. (2024, January 22). *What is an ideal customer profile?* Unusual Ventures – Your Partner For The Journey Ahead. <https://www.unusual.vc/articles/what-is-an-ideal-customer-profile>.

Walker, R. (2017). *What Y combinator looks for in nonprofits.* Y Combinator. <https://www.ycombinator.com/blog/what-y-combinator-looks-for-in-nonprofits/>.

Woo, E., & Afrati, A. (2023, May 4). *OpenAI's losses doubled to \$540 million as it developed ChatGPT.* The Information. <https://www.theinformation.com/articles/openais-losses-doubled-to-540-million-as-it-developed-chatgpt>.

Y Combinator. (2023). *Top companies by valuation.* <https://www.ycombinator.com/topcompanies/valuation/public>.